



MUCHANGI PATRICK & ASSOCIATES

ADVOCATES OF THE HIGH COURT OF KENYA · DIGITAL TRUST ASSURANCE

— REGULATORY ANALYSIS SERIES — NO. 003

Artificial Intelligence Oversight: The Board's Newest *Fiduciary Duty*

Why governing AI can no longer wait for the law that will eventually require it.

PREPARED FOR

Boards & Executive
Leadership

SECTORS

Financial Services · Healthcare ·
Technology

PUBLICATION

Knowledge Centre,
2026

CONTENTS

In This Analysis

Exec.	Executive Summary	<u>3</u>
01	Why AI Governance Belongs at Board Level	<u>3</u>
02	The Shape of Kenya’s Emerging AI Framework	<u>4</u>
03	Where AI Risk Actually Lives Inside an Institution	<u>5</u>
04	The Governance Gap Most Boards Have Right Now	<u>5</u>
05	Building AI Oversight Before Regulation Compels It	<u>6</u>
06	What Boards Should Be Asking Management	<u>6</u>
—	Govern Ahead of the Curve, Not Behind It	<u>7</u>

This analysis forms part of the Muchangi Patrick & Associates Advocates Knowledge Centre — the firm’s Regulatory Analysis Series for Boards and executive leadership in Kenya.

Artificial Intelligence Oversight: The Board's Newest Fiduciary Duty

Why governing AI can no longer wait for the law that will eventually require it.

EXECUTIVE SUMMARY

Artificial intelligence has moved from pilot project to production system inside many Kenyan institutions faster than governance structures have kept pace. Credit scoring, diagnostic support, fraud detection and customer service now routinely run on models that few Boards could describe in detail, let alone audit.

Kenya's regulatory framework for AI is still taking shape, but the absence of a dedicated statute does not mean the absence of exposure. Existing obligations around fairness, accountability, data protection and consumer harm already reach into how AI systems are built, bought and deployed.

Boards that wait for legislation before establishing AI oversight will, by the time it arrives, be retrofitting governance onto systems already embedded in daily operations. The more defensible position is to govern now, on the assumption that regulation will simply formalise a standard already being met.

01

Why AI Governance Belongs at Board Level

An algorithm that scores a loan application, flags a fraudulent transaction, or triages a patient is making decisions that used to sit with a trained employee — decisions that carried accountability, a paper trail, and a person who could explain the reasoning.

When a model makes that decision instead, the accountability does not disappear. It moves, by default, to whoever deployed the system: the institution, and ultimately its Board.

An institution that cannot explain how its AI system reached a decision has not automated a process. It has outsourced its judgment to a system it does not fully understand.

02

The Shape of Kenya's Emerging AI Framework

Kenya's draft AI legislation signals an intent to regulate AI systems according to the risk they pose, echoing the approach taken in the EU's AI Act and discussed in global forums such as the UN's dialogue on AI governance. The details will continue to be negotiated, but the direction of travel — tiered obligations rising with the sensitivity of the use case — is consistent across jurisdictions now legislating in this space.

Institutions using AI in credit decisions, healthcare diagnostics or biometric identification should assume they will eventually sit in a higher-obligation tier, whatever the final statute says, because that is where comparable frameworks elsewhere have placed those use cases.

Where AI Risk Actually Lives Inside an Institution

- Procurement — AI embedded in a vendor's platform, adopted without a governance review
- Automated decisioning — credit scoring, underwriting or diagnostic triage running with limited human oversight
- Generative AI in daily use — staff pasting confidential or client data into public AI tools
- Bias and explainability — models trained on historical data that may encode historical inequities

The Governance Gap Most Boards Have Right Now

Ask most Boards to name every AI system currently in production, including the ones embedded inside vendor platforms, and the honest answer is usually incomplete. That gap — not knowing what is running, rather than knowing and choosing to accept the risk — is the more urgent problem, because it cannot be managed if it is not first inventoried.

Building AI Oversight Before Regulation Compels It

- A current inventory of every AI system in production, including vendor-embedded tools
- A human-in-the-loop requirement for any decision materially affecting a customer or patient
- A written policy governing staff use of generative AI tools, including what may and may not be entered into them
- Vendor due diligence that asks how a model was trained and validated, not only what it costs

What Boards Should Be Asking Management

Every Board should be able to put the following questions to management, and expect a considered answer:

- | | |
|----|---|
| 01 | Can we name every AI system currently in production, including those embedded in vendor platforms? |
| 02 | Who is accountable when an automated decision is wrong? |
| 03 | Is there a human in the loop for decisions that materially affect a customer or patient? |
| 04 | Do we have a policy governing staff use of generative AI tools? |
| 05 | How was each material model trained, and have we validated it ourselves or only trusted the vendor? |
| 06 | What would we do if a regulator asked us to explain a specific automated decision? |
| 07 | Is AI governance reviewed by the Board, or delegated entirely to the technology team? |

CONCLUSION

Govern Ahead of the Curve, Not Behind It

Regulation will eventually catch up to where AI already operates inside Kenyan institutions. Boards that govern ahead of that curve will meet the coming framework as a formality; those that wait will meet it as an audit.

BOARD CONSIDERATION

“Could your Board name every material AI system currently making decisions about your customers or patients?”

If the honest answer is “not fully,” that inventory — not the eventual statute — is where oversight should begin.

M

About the Author

Muchangi Patrick & Associates Advocates is a Kenyan law firm specialising in Digital Trust Assurance, data protection, technology law, AI governance and regulatory advisory. We help Boards understand, govern and continuously improve how their organisations manage information, technology, artificial intelligence and legal risk.

Request an Executive Briefing

We help leadership teams understand, govern and continuously improve Digital Trust through structured, evidence-based advisory engagements.

Pinetree Plaza, Nairobi · 0722 878 607 · consult@dataprivacyadvocates.co.ke